

COMP/EECE 7/8740 Neural Networks

Topics:

- CNN Architectures
 - Res2Net
 - Hybrid Networks
 - FactralNet and DenseNet
 - MobileNet and so on

Md Zahangir Alom
Department of Computer Science
University of Memphis, TN

What are the current trends?

- Recent advances in backbone convolutional neural networks (CNNs) continually demonstrate stronger **multi-scale representation ability**
- Leading to **consistent performance gains** in a wide range of applications.
- Current trends:
 - Minor modifications to ResNets
 - Biggest trend is to split of into several branches, and merge through summation
 - A couple architectures go crazy with branch & merge, without explicit identity connections

What are the current trends?

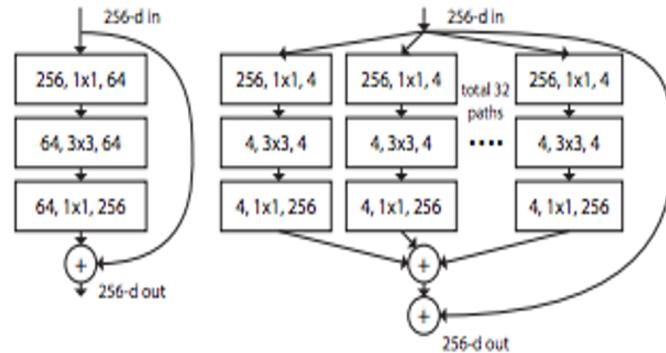


Figure 1. Left: A block of ResNet [13]. Right: A block of ResNeXt with cardinality = 32, with roughly the same complexity. A layer is shown as (# in channels, filter size, # out channels).

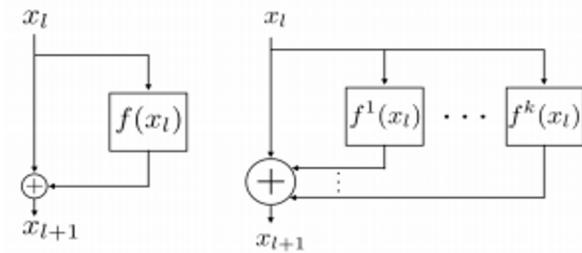


Figure 2: A residual block (left) versus a multi-residual block (right).

ResNeXt

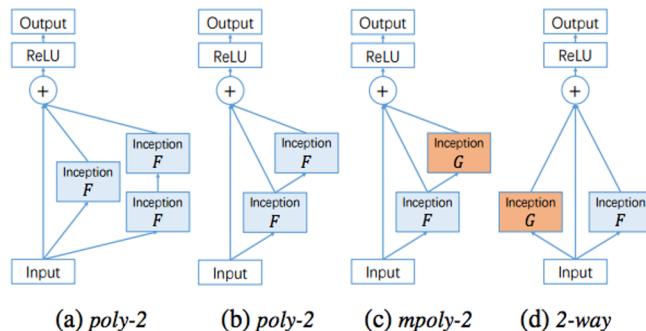
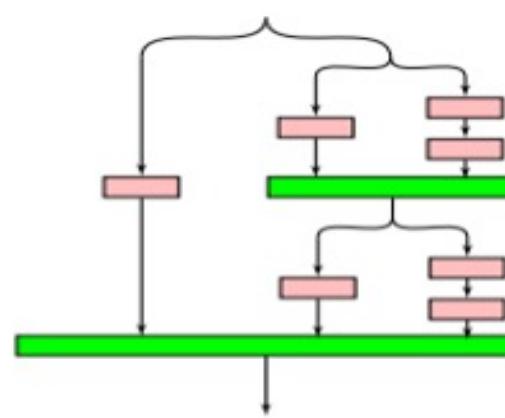


Figure 4: Examples of PolyInception structures.

PolyNet

MultiResNet



FactalNet

Aggregated Residual Transformations

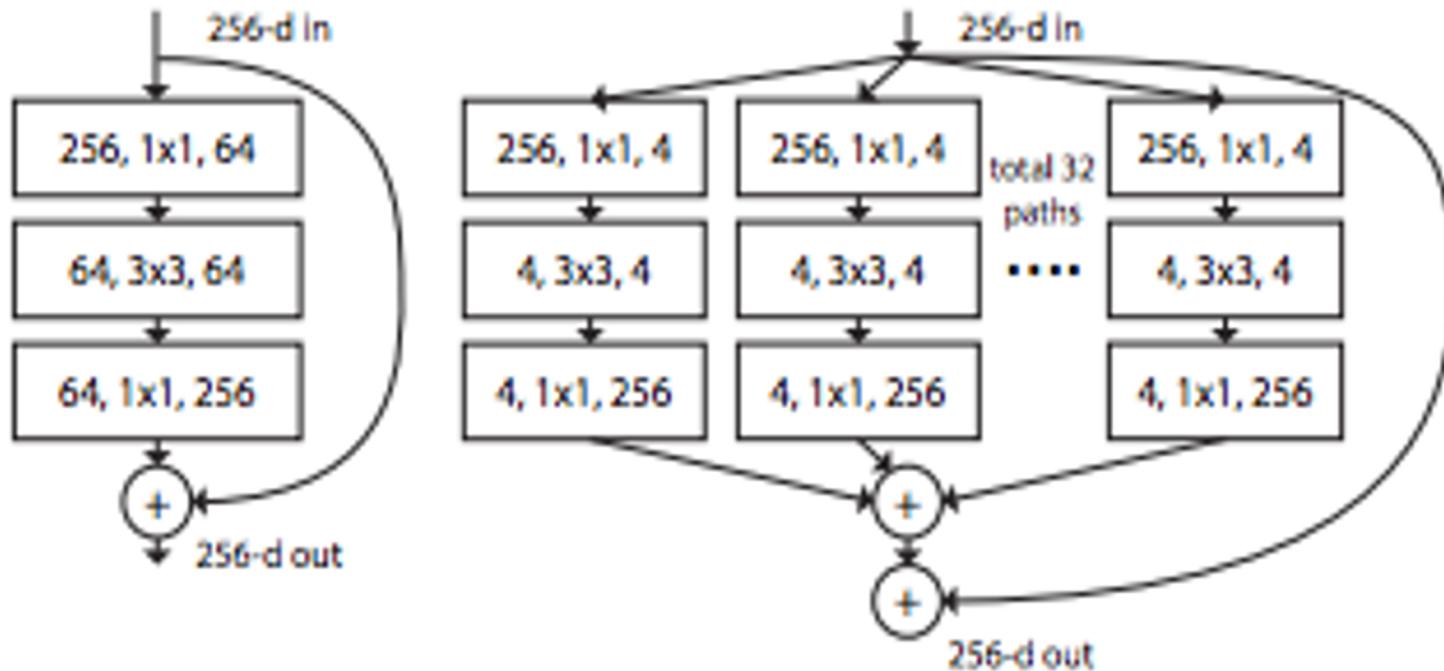
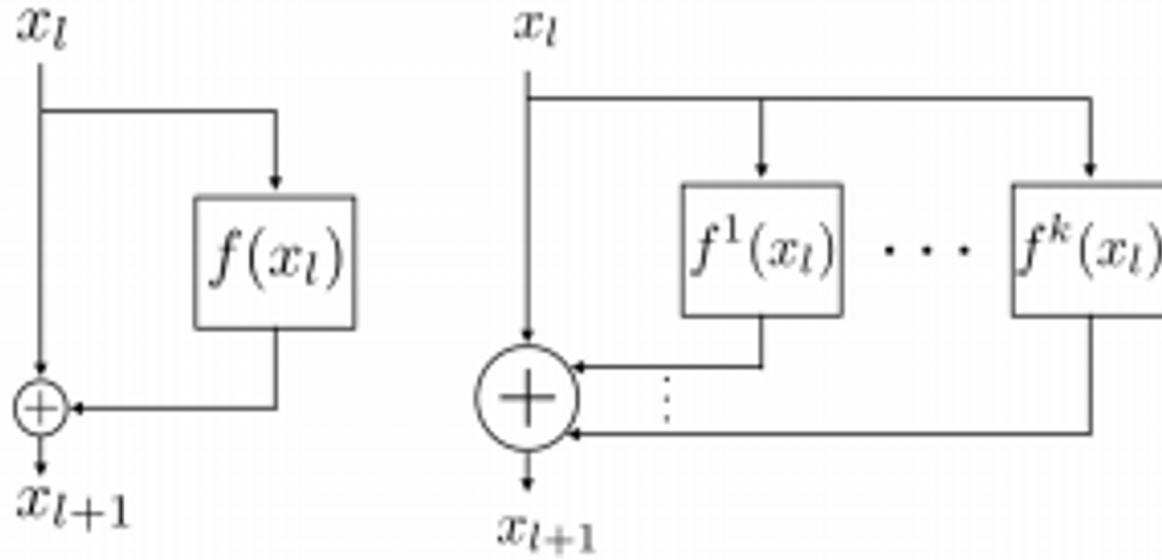


Figure 1. **Left:** A block of ResNet [13]. **Right:** A block of ResNeXt with cardinality = 32, with roughly the same complexity. A layer is shown as (# in channels, filter size, # out channels).

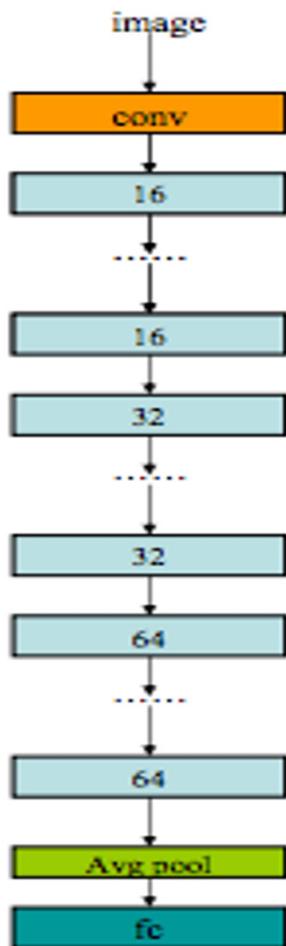
Multi-Residual Networks (ResNet)



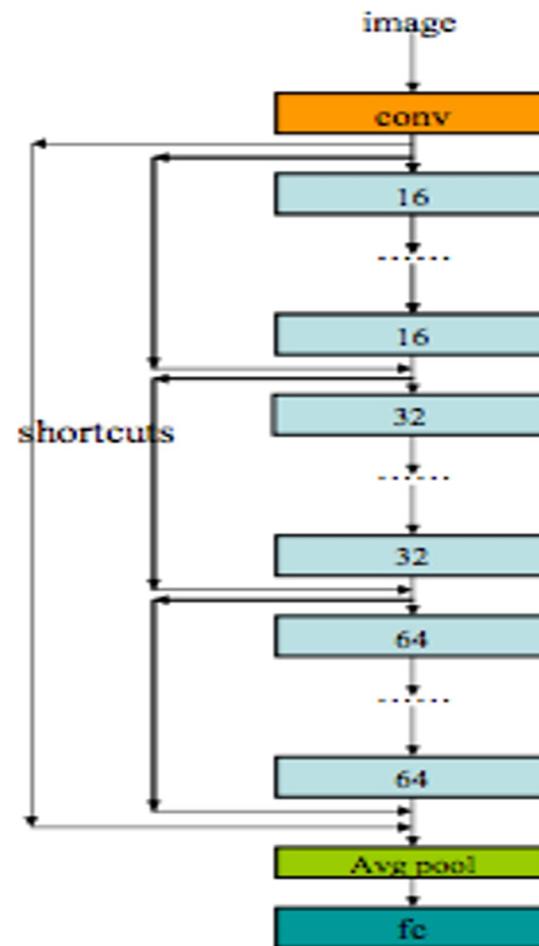
Multi-Residual Networks (ResNet) : residual block (left)
versus a multi-residual block (right)

Residuals of Residual: try going meta

- Residuals of Residuals:



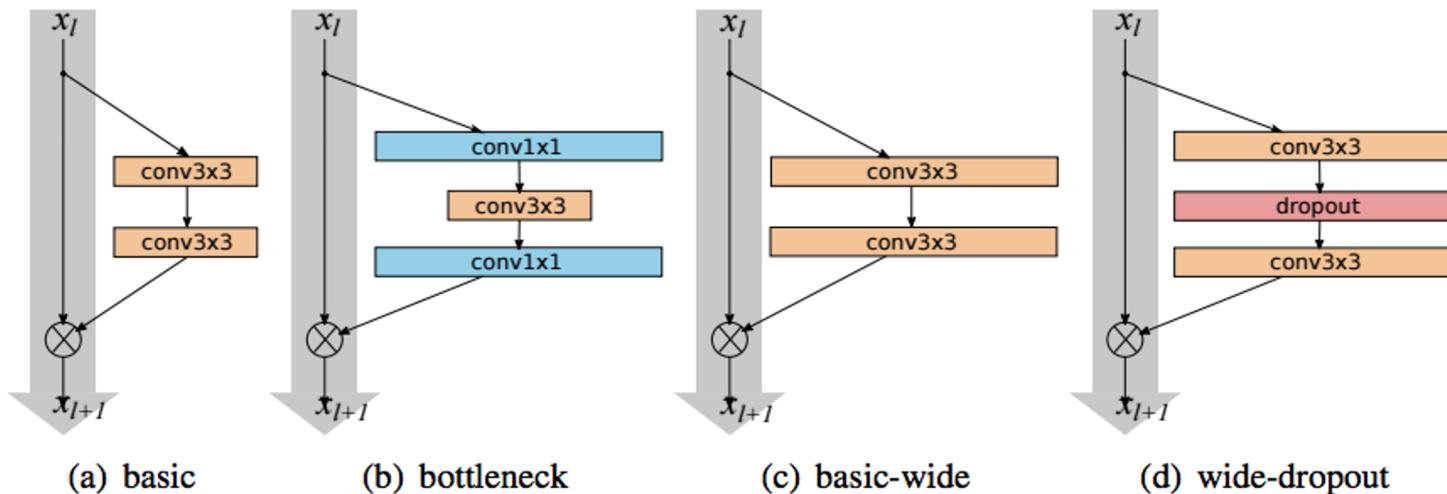
Residual Networks



RoR

ResNet tweaks: Wide ResNets

- Use pre-activation ResNet's basic block with more feature maps
- Used parameter “ k ” to encode width
- Investigated relationship between width and depth to find a good tradeoff



Various residual blocks used in the paper. Batch normalization and ReLU precede each convolution (omitted for clarity)

ResNet tweaks: Wide ResNets

- Use pre-activation ResNet's basic block with more feature maps
- **Used parameter “k” to encode width**
- Investigated relationship between width and depth to find a good tradeoff

group name	output size	block type = $B(3, 3)$
conv1	32×32	$[3 \times 3, 16]$
conv2	32×32	$\begin{bmatrix} 3 \times 3, 16 \times k \\ 3 \times 3, 16 \times k \end{bmatrix} \times N$
conv3	16×16	$\begin{bmatrix} 3 \times 3, 32 \times k \\ 3 \times 3, 32 \times k \end{bmatrix} \times N$
conv4	8×8	$\begin{bmatrix} 3 \times 3, 64 \times k \\ 3 \times 3, 64 \times k \end{bmatrix} \times N$
avg-pool	1×1	$[8 \times 8]$

Structure of wide residual networks. ResNet block of type B(3,3). Network width is determined by factor k . Groups of convolutions are shown in brackets where N is a number of blocks in group

ResNet tweaks: Wide ResNets

- Use pre-activation ResNet's basic block with more feature maps
- Used parameter “ k ” to encode width
- **Investigated relationship between width and depth to find a good tradeoff**

depth	k	# params	CIFAR-10	CIFAR-100
40	1	0.6M	6.85	30.89
40	2	2.2M	5.33	26.04
40	4	8.9M	4.97	22.89
40	8	35.7M	4.66	-
28	10	36.5M	4.17	20.50
28	12	52.5M	4.33	20.43
22	8	17.2M	4.38	21.22
22	10	26.8M	4.44	20.75
16	8	11.0M	4.81	22.07
16	10	17.1M	4.56	21.59

Test error (%) of various wide networks on CIFAR-10 and CIFAR-100 (ZCA preprocessing).

ResNet tweaks: Wide ResNets

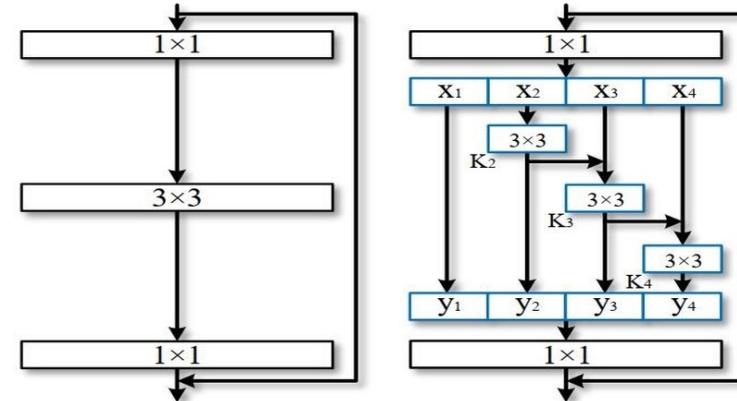
- These obtained state of the art results on CIFAR datasets
- Were outperformed by bottlenecked networks on ImageNet
- **Best results on ImageNet were obtained by widening ResNet-50**

Model	top-1 err, %	top-5 err, %	#params	time/batch 16
ResNet-50	24.01	7.02	25.6M	49
ResNet-101	22.44	6.21	44.5M	82
ResNet-152	22.16	6.16	60.2M	115
WRN-50-2-bottleneck	21.9	6.03	68.9M	93
pre-ResNet-200	21.66	5.79	64.7M	154

“With widening factor of 2.0 the resulting WRN-50-2-bottleneck outperforms ResNet- 152 having 3 times less layers, and being significantly faster.”

Res2Net: A New Multi-scale Backbone Architecture

- Most existing methods represent the multi-scale features in a layer-wise manner

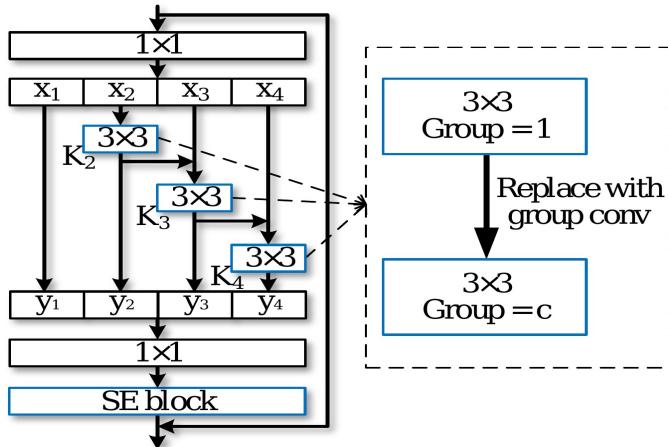


(a) Bottleneck

(b) Res2Net module

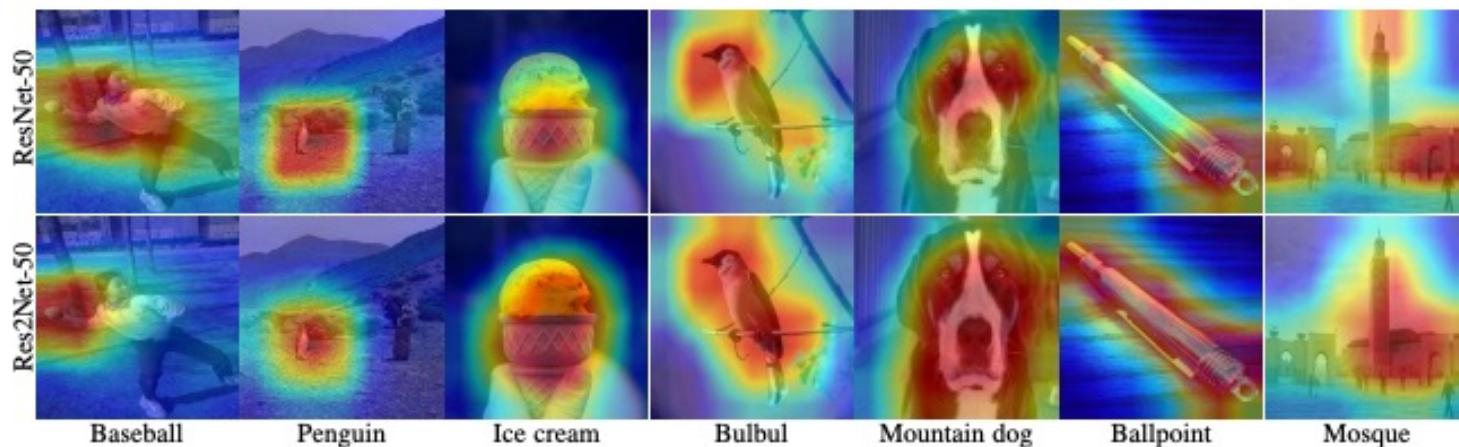
- A novel building block for CNNs, namely **Res2Net**, by constructing **hierarchical residual-like connections within one single residual block**.
- The Res2Net represents multi-scale features at a **granular level and increases the range of receptive fields** for each network layer

Res2Net: A New Multi-scale Backbone Architecture



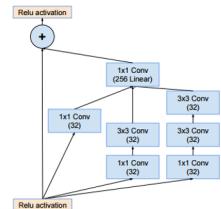
The Res2Net module can be integrated with the dimension cardinality (replace conv with group conv) and SE blocks.

- The proposed Res2Net block can be plugged into the state-of-the-art backbone CNN models, e.g., ResNet, ResNeXt, and DLA

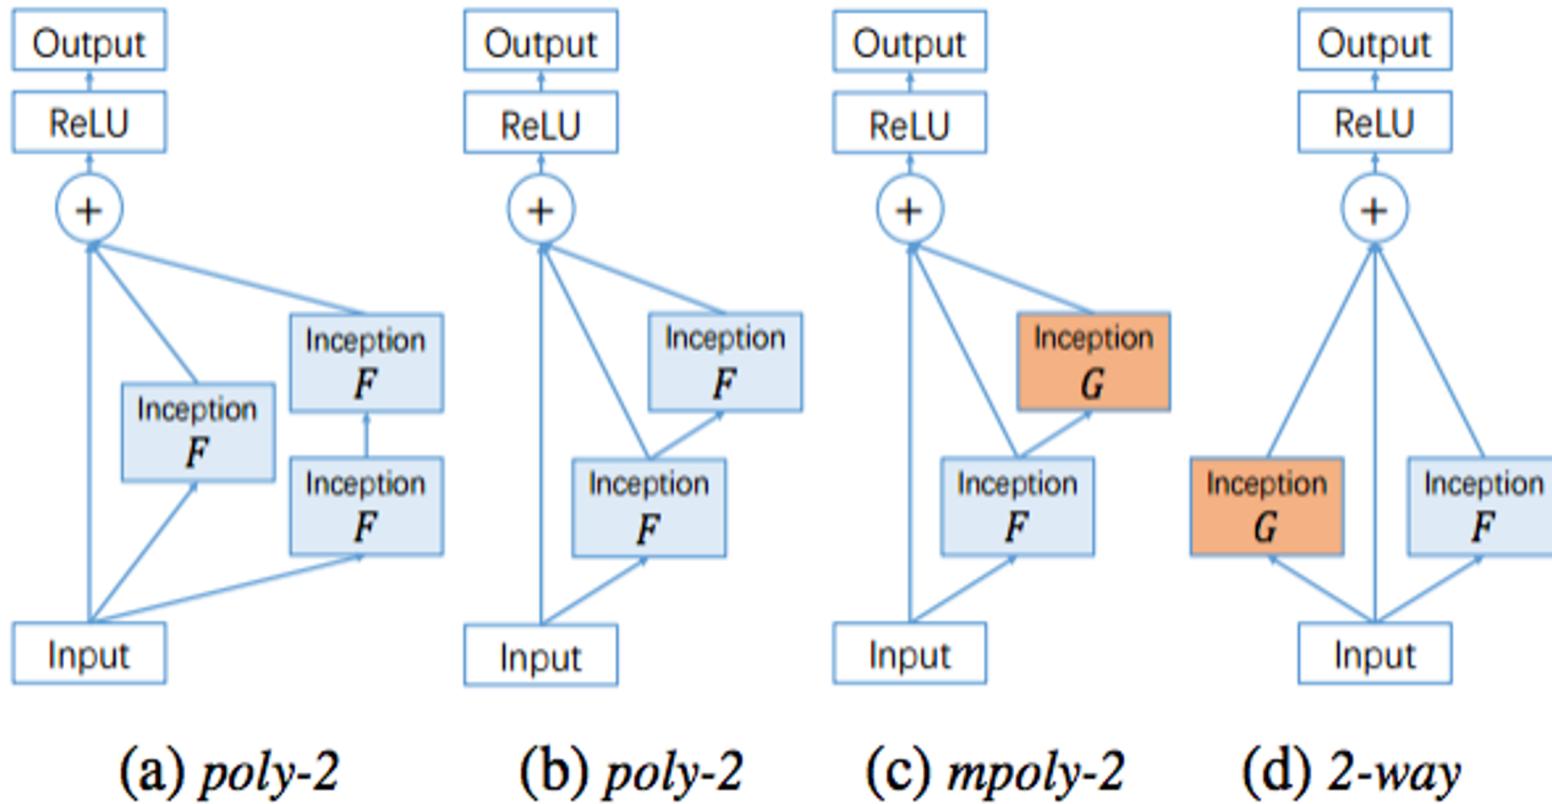


Visualization of class activation mapping, using ResNet-50 and Res2Net-50 as backbone networks.

PolyNet



Inception ResNet



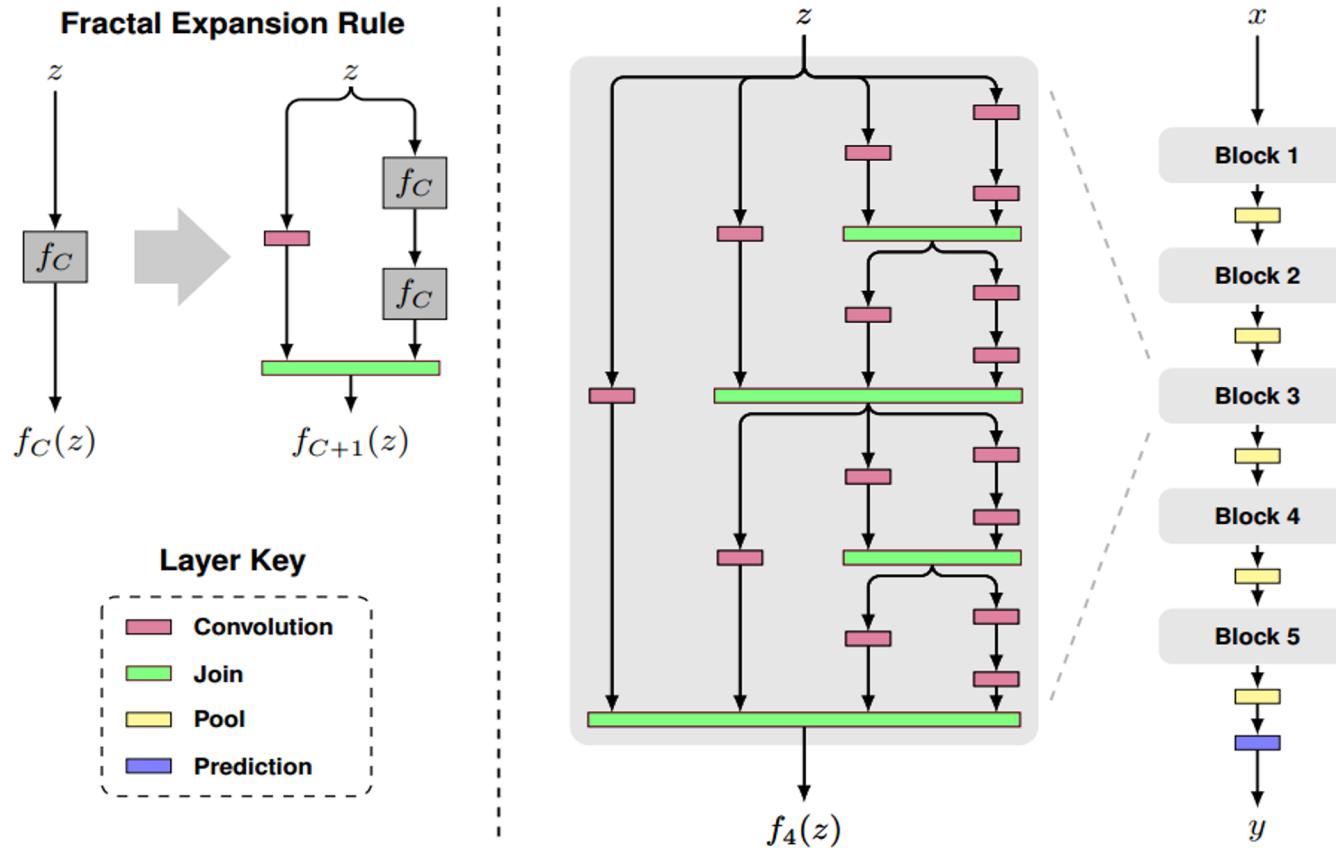
PolyNet Structures

Aside from ResNets, Inception, and Inception-ResNet

FractalNet and DenseNet

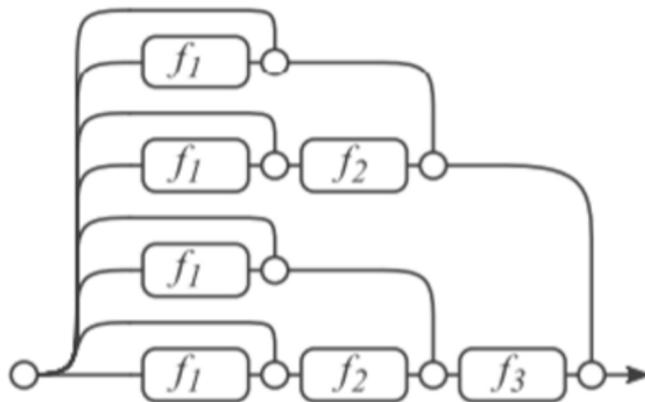
FractalNet

- A extremely deep architecture that does not rely on residuals

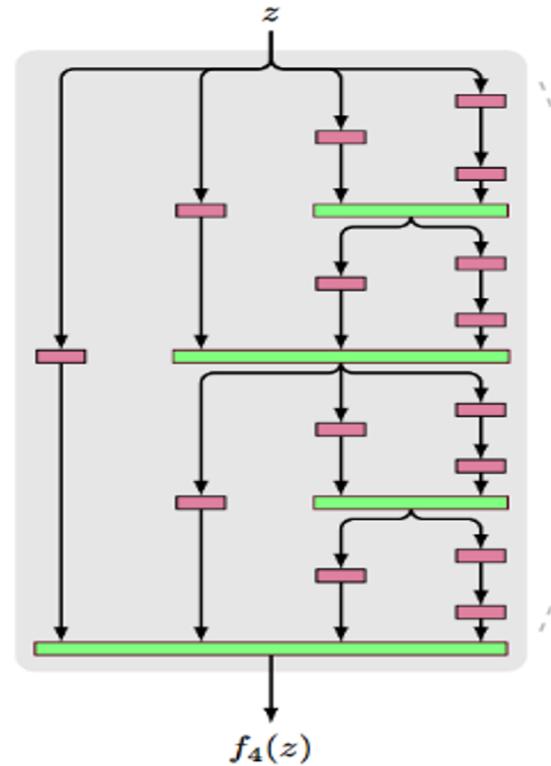


FractalNet

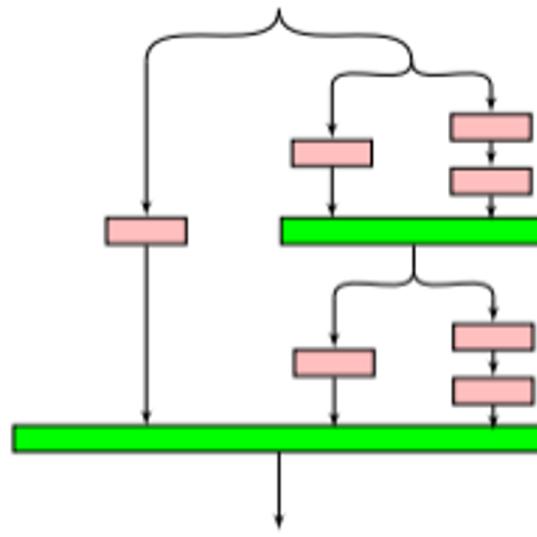
- A extremely deep architecture that **does not rely on residuals**
- Interestingly, its architecture is **similar to an unfolded ResNet**



(b) Unraveled view of (a)

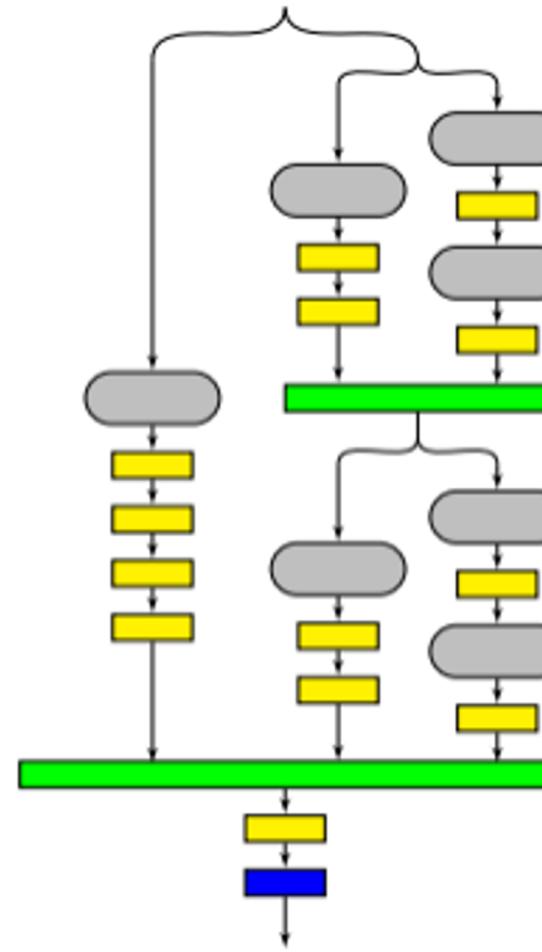


Fractal of Fractals : try going meta



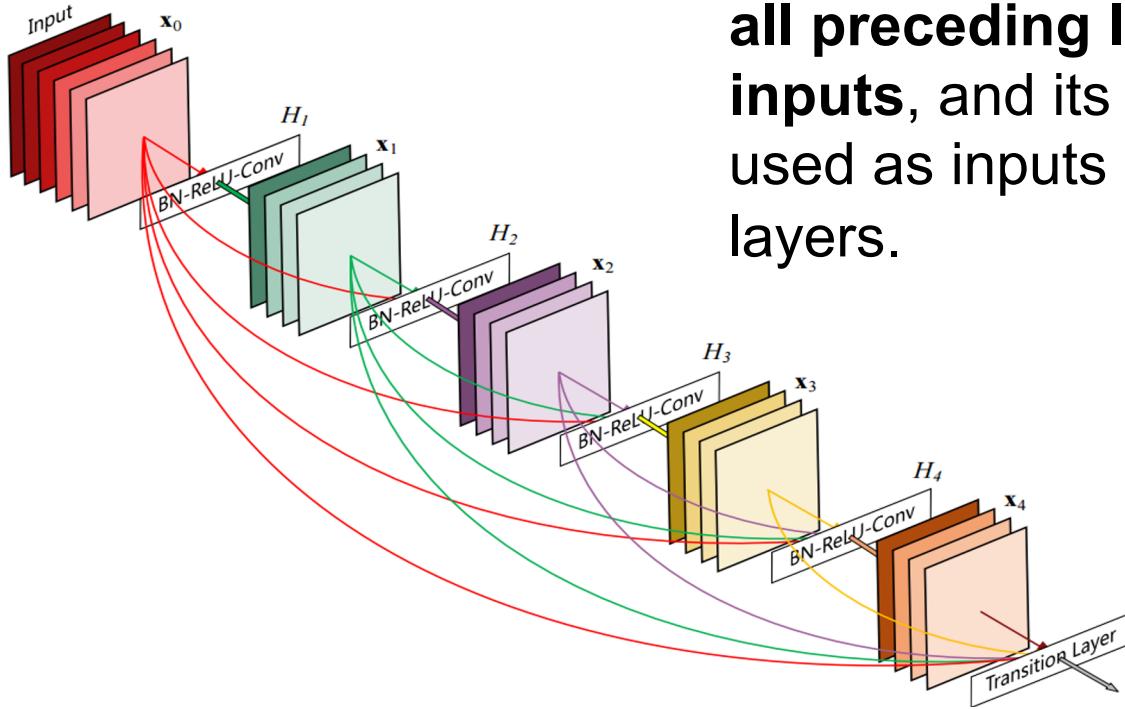
- [pink rectangle] convolutional layer
- [yellow rectangle] pooling layer
- [blue rectangle] prediction layer
- [green bar] joining layer
- [grey oval] FractalNet module

(a) FractalNet module



(b) Fractal of FractalNet architecture

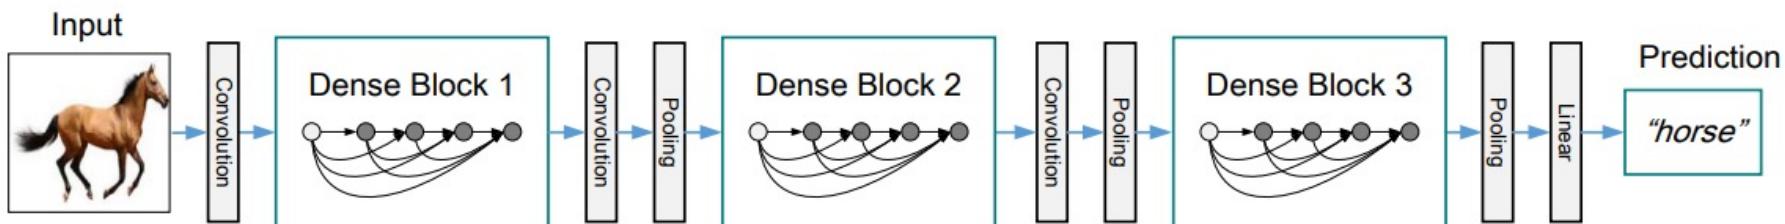
DenseNet (Within a DenseBlock)



- Every layer is connected to all other layers.
- For each layer, the **feature-maps of all preceding layers are used as inputs**, and its own feature-maps are used as inputs into all subsequent layers.

DenseNet

- Consists with **multiple dense and transition blocks**
- Basic operations:
 - Convolutions (1x1,3x3)
 - Batch Normalization (BN)
 - ReLU and
 - Pooling.



Overall Network Architecture

Advantages of DenseNet

- Alleviate the **vanishing-gradient problem**
- Strengthen of **feature propagation**
- Encourage **feature reuse**
- **Substantially** reduce the number of the network parameters.

MobileNets

- Introduce two simple global hyperparameters that efficiently trade off between latency and accuracy
- These hyper-parameters allow the model builder to choose the right sized model for their application based on the constraints of the problem

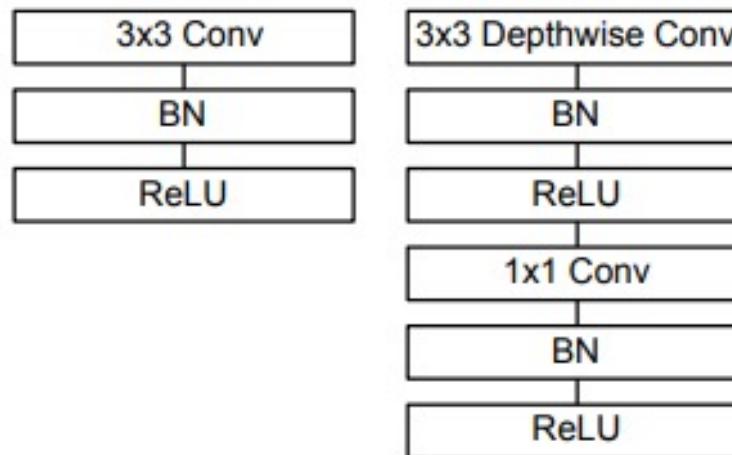
MobileNets



MobileNet models can be applied to various recognition tasks for efficient on device intelligence

MobileNets

- The MobileNet model is based on depthwise separable convolutions a **1×1 convolution called a pointwise convolution.**
 - The Depthwise convolution applies a single filter to each input channel.
 - The pointwise convolution then applies a 1×1 convolution to combine the outputs the depthwise convolution



Left: Standard convolutional layer with batchnorm and ReLU. Right: Depthwise Separable convolutions with Depthwise and Pointwise layers followed by batchnorm and ReLU.

MobileNets Architecture

Table 1. MobileNet Body Architecture

Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32$ dw	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64$ dw	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128$ dw	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256$ dw	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
5× Conv dw / s1	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
	$1 \times 1 \times 512 \times 512$	$14 \times 14 \times 512$
Conv dw / s2	$3 \times 3 \times 512$ dw	$14 \times 14 \times 512$
Conv / s1	$1 \times 1 \times 512 \times 1024$	$7 \times 7 \times 512$
Conv dw / s2	$3 \times 3 \times 1024$ dw	$7 \times 7 \times 1024$
Conv / s1	$1 \times 1 \times 1024 \times 1024$	$7 \times 7 \times 1024$
Avg Pool / s1	Pool 7×7	$7 \times 7 \times 1024$
FC / s1	1024×1000	$1 \times 1 \times 1024$
Softmax / s1	Classifier	$1 \times 1 \times 1000$

- Significantly reduce the network parameters and computational power.

Table 2. Resource Per Layer Type

Type	Mult-Adds	Parameters
Conv 1×1	94.86%	74.59%
Conv DW 3×3	3.06%	1.06%
Conv 3×3	1.19%	0.02%
Fully Connected	0.18%	24.33%

Bonus Material!

We'll cover spatial transformer networks (briefly)

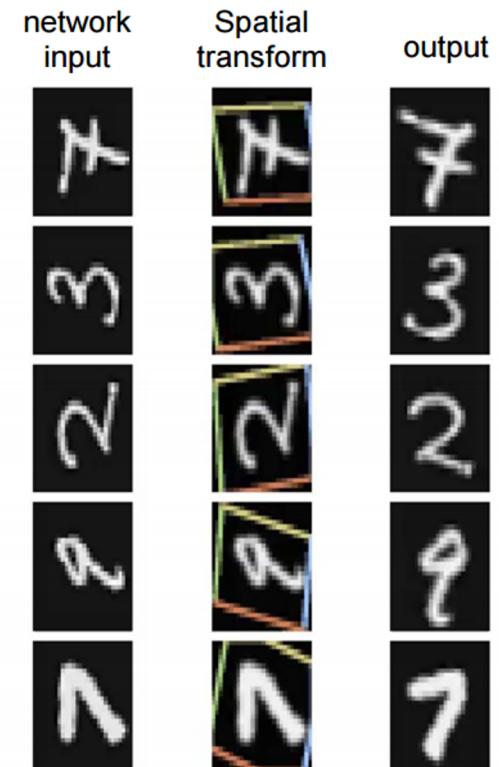
Spatial Transformer Networks

A module to provide spatial transformation capabilities on individual data samples.

Idea:

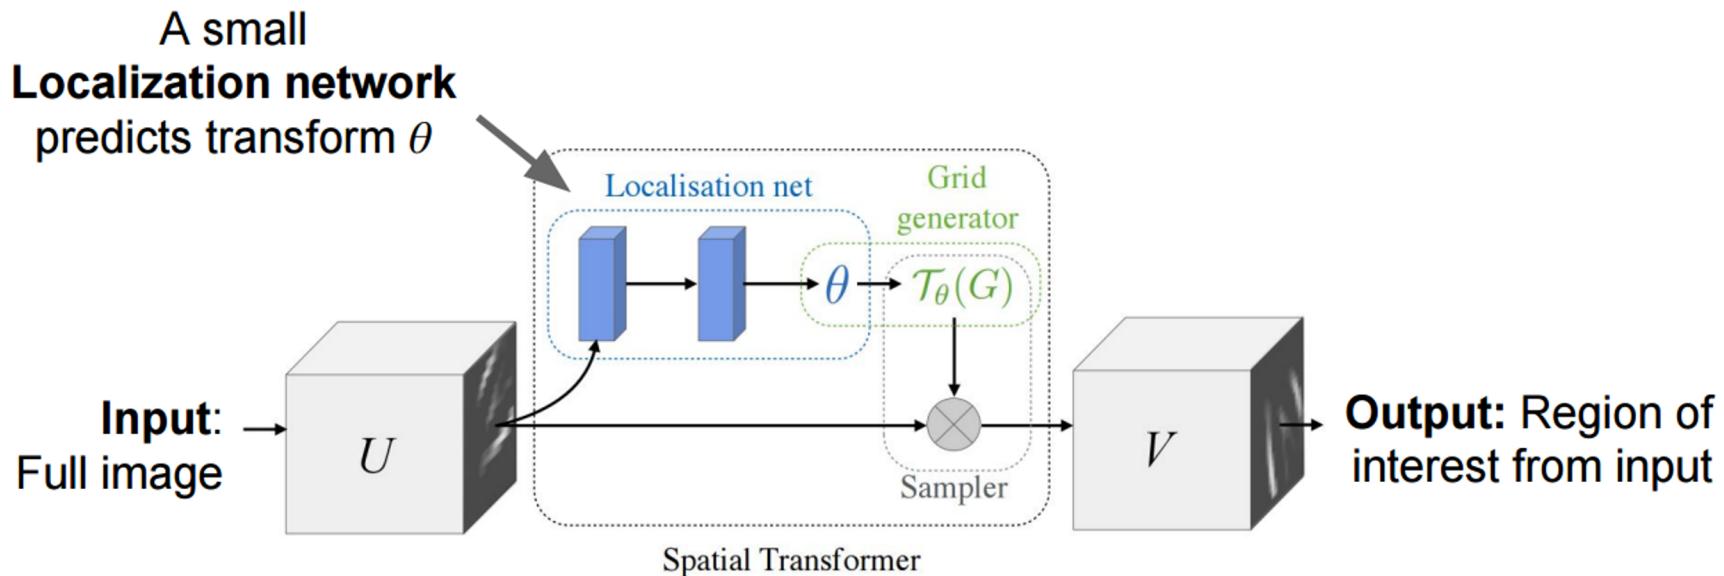
Function mapping pixel coordinates of output to pixel coordinates of input.

$$\begin{pmatrix} x_i^s \\ y_i^s \end{pmatrix} = \begin{bmatrix} \theta_{11} & \theta_{12} & \overleftarrow{\theta_{13}} \\ \theta_{21} & \theta_{22} & \theta_{23} \end{bmatrix} \begin{pmatrix} x_i^t \\ y_i^t \\ 1 \end{pmatrix}$$

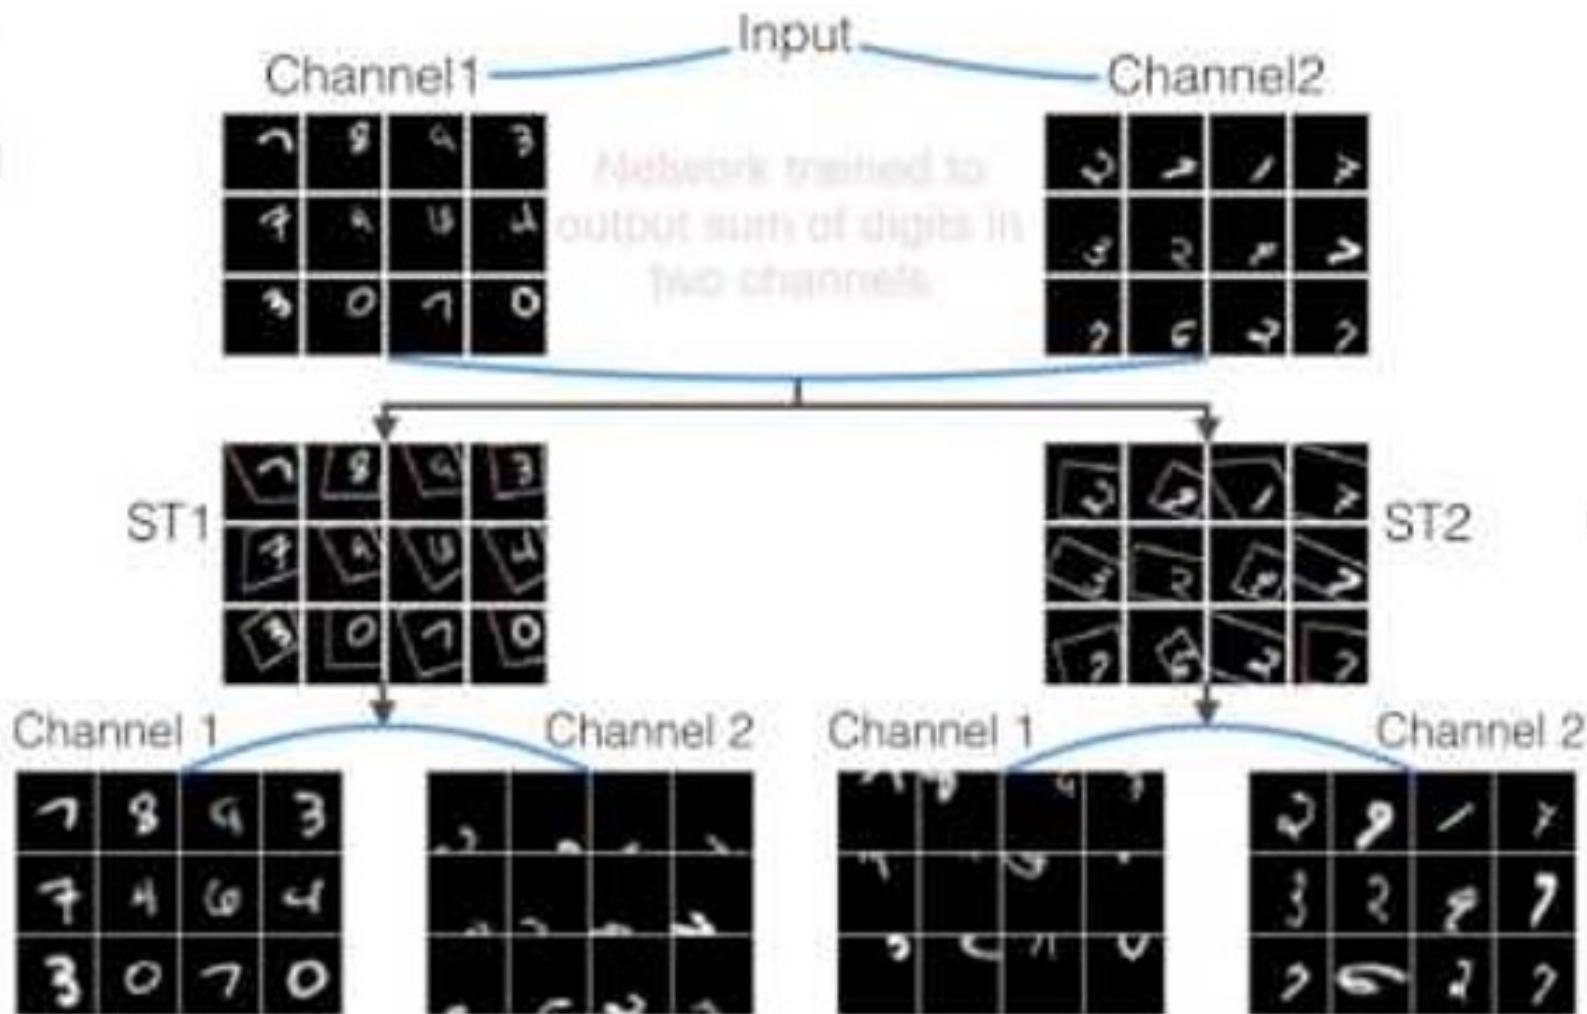


Spatial transform by how much?

The localisation network function can take any form, such as a **fully-connected network** or a **convolutional network**, but should include a final **regression layer** to produce the **transformation parameters θ** .



MNIST Addition



Recap (General model design principles)

- At surface level, there's tons of new architectures that are very different
- Upon closer inspection, **most of them are reapplying well established principles**
- Universal principles seem to be **having shorter subpaths through the networks**
- Use **skip connections and/or create multiple paths** through the network, identity propagation (Residuals, Dense Blocks) seem to make training easier
- **Reduce filter sizes** (except possibly at the lowest layer), factorize filters aggressively
- Use **1x1 convolutions to reduce and expand** the number of feature maps judiciously

What's missing here?

- Data augmentation for training
- Training tricks and details:
 - initialization
 - Batch Normalization
 - Regularization : Dropout, drop-connection
 - Advanced activation functions
 - Better optimization function: SGD, AdaDelta, AdaGrad.....
- Alternative of fully-connected layers
- Network compression and parameter optimization

References

- Sergey Ioffe, Christian Szegedy, Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift, <http://jmlr.org/proceedings/papers/v37/ioffe15.pdf>
- K. He, X. Zhang, S. Ren, and J. Sun, Deep Residual Learning for Image Recognition, CVPR 2016 <https://arxiv.org/abs/1512.03385>
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun, Identity Mappings in Deep Residual Networks <https://arxiv.org/abs/1603.05027>
- Saining Xie, Ross Girshick, Piotr Dollár, Zhuowen Tu, Kaiming He, Aggregated Residual Transformations for Deep Neural Networks, <https://arxiv.org/pdf/1611.05431v1.pdf>
- Max Jaderberg, Karen Simonyan, Andrew Zisserman, Koray Kavukcuoglu: Spatial Transformer Networks <https://arxiv.org/abs/1506.02025>
- Leslie N. Smith, Nicholay Topin, Deep Convolutional Neural Network Design Patterns, <https://arxiv.org/abs/1611.00847>
- Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, Alex Alemi: Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning <https://arxiv.org/abs/1602.07261>
- Gustav Larsson, Michael Maire, Gregory Shakhnarovich, FractalNet: Ultra-Deep Neural Networks without Residuals <https://arxiv.org/abs/1605.07648>
- Gao Huang, Zhuang Liu, Kilian Q. Weinberger, Laurens van der Maaten: Densely Connected Convolutional Networks <https://arxiv.org/pdf/1608.06993v3.pdf>
- Rupesh Kumar Srivastava, Klaus Greff, Jürgen Schmidhuber: Highway Networks <https://arxiv.org/abs/1505.00387>
- Xingcheng Zhang, Zhizhong Li, Chen Change Loy, Dahua Lin, PolyNet: A Pursuit of Structural Diversity in Very Deep Networks, <https://arxiv.org/abs/1611.05725>
- Jost Tobias Springenberg, Alexey Dosovitskiy, Thomas Brox, and Martin Riedmiller. Striving for simplicity: The all convolutional net. <https://arxiv.org/abs/1412.6806>
- Andreas Veit, Michael Wilber, Serge Belongie, Residual Networks Behave Like Ensembles of Relatively Shallow Networks, <https://arxiv.org/pdf/1605.06431v2.pdf>
- Klaus Greff, Rupesh K. Srivastava & Jürgen Schmidhuber, Highway and Residual Networks Learn Unrolled Iterative Estimation, <https://arxiv.org/pdf/1612.07771v1.pdf>
- Min Lin, Qiang Chen, Shuicheng Yan, Network In Network, <https://arxiv.org/abs/1312.4400>
- Brian Chu, Daylen Yang, Ravi Tadinada, Visualizing Residual Networks, <https://arxiv.org/abs/1701.02362>
- Djork-Arné Clevert, Thomas Unterthiner, Sepp Hochreiter: Fast and Accurate Deep Network Learning by Exponential Linear Units (ELUs) <https://arxiv.org/abs/1511.07289>
- Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun, Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification, <https://arxiv.org/abs/1502.01852>
- Anish Shah, Eashan Kadam, Hena Shah, Sameer Shinde, Sandip Shingade, Deep Residual Networks with Exponential Linear Unit, <https://arxiv.org/abs/1604.04112>